

Development of a Search Engine and Applications in IR

Robert Pinsler
WKWSCI, NTU
N1509281G@e.ntu.edu.sg

Yiding Liu
SCE, NTU
LIUY0130@e.ntu.edu.sg

Yitong Guan
SCE, NTU
GUAN0049@e.ntu.edu.sg

Jenn Bing Ong
SCE, NTU
ONGJ0063@e.ntu.edu.sg

ABSTRACT

1. INTRODUCTION

2. SYSTEM OVERVIEW

The search engine is able to index and search publication records listed in the dblp computer science bibliography (DBLP) [2]. It is built on top of PyLucene, a Python extension of the open-source software library Lucene¹ that provides text indexing and search capabilities.

Publication records from DBLP are provided in a single XML file. The dataset comprises of various kinds of record types, such as articles published in a journal or magazine (*article*), papers published in a conference or workshop proceedings (*inproceedings*), books, and PhD theses. From those, only documents that are classified as *article* or *inproceedings* are considered.

3. INDEXING

To index a record, it first has to be extracted from the XML file before it can be stored. This can be done conveniently using an XML parser. Due to the large file size of the dataset, reading in the whole XML tree at once is impractical. Instead, we use a SAX-like parser from the lxml² Python library that sequentially reads the document and emits events when it encounters certain XML tags. This allows to only react to `<article>` and `<inproceedings>` tags, thereby ignoring record types we are not interested in. This makes it very fast and memory-efficient. During event processing, it is possible to access the current element, i.e. an *article* or *inproceedings* record, and its children, i.e. attributes of the record such as title or year of publication. As some of the publication titles use basic HTML text formatting, HTML tags are removed prior to indexing.

Each record is indexed as a Lucene *document*, comprising of the following *fields*: id, title, authors, year and venue

(journal or booktitle, depending on the record type). To facilitate the search, an additional *content* field is added that concatenates the values of the above fields except the id. All fields are treated as strings³. Contents of the original fields (i.e. without the *content* field) are stored in the index, which allows to retrieve them later during search. This is especially useful to enhance search results with additional information when presented to the user. The values of the fields *id* and *year* are stored as-is. All other fields are tokenized (f_{token}) and further processed, including lowercase conversion (f_{low}), removal of stopwords (f_{stop}) and stemming (f_{stem}). In particular, we use the Porter stemmer and the default stopword list of Lucene, which consists of 33 common English words. Table 1 summarizes the analysis applied to each field.

Table 1: Analysis applied to document fields

Field	f_{token}	f_{low}	f_{stop}	f_{stem}
id				
year				
authors	X	X	X	
venue	X	X	X	
content	X	X	X	
title	X	X	X	X

In order to evaluate the impact of additional token processing onto the index, different configurations of lowercase conversion, removal of stopwords and/or stemming are applied to the *title* field. We measure the speed of indexing (including parsing and analyzing the documents) in seconds, t_{ind} , as well as the number of terms in the vocabulary, $|V_{\text{title}}|$. The results are shown in Table 2. Note that indexing speed slightly varies over different runs.

Table 2: Evaluation of token processing on *title* field

f_{low}	f_{stop}	f_{stem}	t_{ind}	$ V_{\text{title}} $
			417.97 (100%)	436 354 (100%)
X			430.57 (+3%)	346 286 (−21%)
	X		414.89 (−1%)	436 321 (±0%)
		X	439.75 (+5%)	365 534 (−16%)
X	X	X	417.94 (±0%)	288 710 (−34%)

Applying lowercase conversion or stemming reduces the vocabulary significantly, whereas the removal of stopwords

³The *year* attribute may also be stored as integer, allowing for features like range searches. However, this would require additional effort during search that we try to avoid.

has nearly no impact. This is expected, since the first two techniques map multiple token onto the same term whereas the exclusion of stopwords decreases the vocabulary size only by the number of stopwords. By combining multiple processing steps, the vocabulary can be further reduced, leading to a reduction to up two thirds of the original size when no further processing is applied. In terms of indexing speed, we found that the additional processing steps are negligible. In particular, there seems to be no positive correlation between indexing speed and the number of processing steps that are applied. One reason for this is that while processing each document is more time-consuming, the cost for indexing a processed document is actually reduced. Therefore, we conclude that, under the given evaluation criteria, it is recommended to apply all three processing techniques due to the reduction in vocabulary size. However, it should be noted that this also has an effect on other properties of the system, e.g. the quality of search results, which might or might not be desired.

4. SEARCH

The search module makes use of the document index to retrieve the most relevant documents to a given query. It supports free text keyword queries on any combination of the attributes title, authors, year and venue. By default, all attributes are considered. This is achieved by utilizing the *content* field of the index. We refer to this as a standard search. Additionally, it is possible to query a combination of specific fields through an advanced search. Those keywords are directly matched against the respective fields in the index. In this case, documents must contain the provided keywords to be returned in the result set. Standard and advanced search can be freely combined during a single search request.

Phrase queries are supported using double quotation marks (e.g. “information retrieval”) in both standard and advanced search. Documents containing a particular phrase receive a higher weight during scoring. This requires an exact match up to the processing on the respective field. If a document does not contain the phrase, it may still be considered when there is a match for at least one of the keywords within the phrase. This considerably increases recall while risking an increase of false positives. Note that there is no restriction for the number of phrases within a query.

The search returns the N most relevant documents along with their ranks, scores, ids and snippets, where N is a configurable parameter. Relevance is determined based on Lucene’s internal scoring function. It also measures the time needed to retrieve the results. The complete procedure for handling queries is outlined in Listing 1.

Listing 1: Query handling

```
def search(self, query, adv_query=None, N=0):
    query = query.strip()
    bq = BooleanQuery()
    if query != '':
        # handle phrases in standard search
        pq, query = self.get_pq(query, 'content')
        if pq is not None:
            bq.add(pq, BooleanClause.Occur.MUST)
    if query != '':
        # handle remaining keywords
        qparser = QueryParser(self.VERSION,
                              'content',
                              self.analyzer)
```

```
        q = qparser.parse(query)
        bq.add(q, BooleanClause.Occur.SHOULD)
    if adv_query is not None:
        for field, query in adv_query.iteritems():
            # handle phrases in advanced search
            pq, query = self.get_pq(query, field)
            if pq is not None:
                bq.add(pq, BooleanClause.Occur.MUST)
            if query != '':
                # handle remaining keywords
                qparser = QueryParser(self.VERSION,
                                      field,
                                      self.analyzer)
                q = qparser.parse(query)
                bq.add(q, BooleanClause.Occur.MUST)
    start = time.clock()
    docs = self.searcher.search(bq, N).scoreDocs
    end = time.clock()
    duration = end - start
    return docs, duration

def get_pq(self, q, field, slop=0, boost=5):
    phrases = re.findall(r'"([^\"]*)"', q)
    if len(phrases) == 0:
        return None, q

    q = re.sub(r'"([^\"]*)"', '', q).strip()
    bq = BooleanQuery()
    for phrase in phrases:
        qparser = QueryParser(self.VERSION,
                              field,
                              self.analyzer)
        # handle phrase and single keywords
        pq = qparser.parse('%s %s"%d^%.1f' %
                           (phrase, phrase,
                            slop, boost))
        bq.add(pq, BooleanClause.Occur.MUST)
    return bq, q
```

5. EVALUATION

6. USER INTERFACE

There are two ways to interact with the search engine: via a text-based command line interface or via a web UI. The latter is built with Flask⁴, a lightweight Python web framework that leverages Jinja2⁵ as a templating engine. This allows to easily create HTML documents from within Python. Additionally, we incorporate the web framework Bootstrap⁶ to achieve responsive web design. Figure 1 depicts the search results for an example query.

7. APPLICATIONS IN IR

We implement two applications based on dblp data. The first one is to find the top-10 most popular research topics of a specific year. The second one is finding top-10 similar publication venues and years with a given publication venue and year. Only paper titles are used in both application.

In this section, the related techniques and the detailed implementations of the applications are demonstrated, followed by the experiment results.

⁴<http://flask.pocoo.org>

⁵<http://jinja.pocoo.org>

⁶<http://getbootstrap.com>

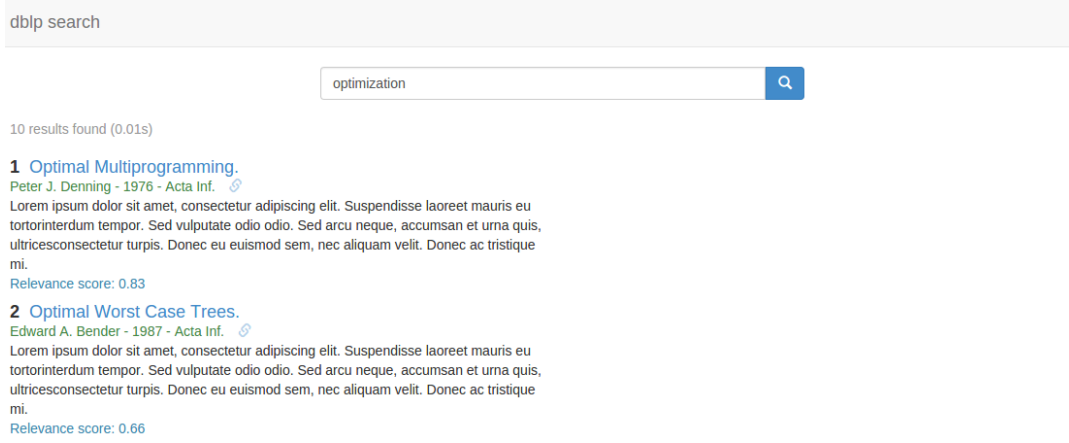


Figure 1: Web-based search user interface

7.1 Popular research topics

Before mining popular research topics from paper titles, we need to firstly extract topics from such data. In our application, we identify meaningful topics from titles based on the patterns of words. More precisely, given a query (e.g. 2014), we have following process to extract topics from it:

- Retrieve the titles of all the papers published in 2014, using the index.
- Tokenize each title and remove punctuations and stop words.
- Use Part-Of-Speech Tagger (POS Tagger)⁷ to assign parts of speech to each word, such as noun, verb, adjective, etc.
- Extract topics from the tagged words based on the given pattern.

From the observation of real topic names, we can conclude that a topic name usually consist of adjectives and nouns. Therefore, in the third step, we define the pattern as $\text{adj}+\text{adj}+\dots+\text{noun}+\text{noun}$. After the topic identification, we count the frequency of each topic appearing in the titles. Finally, the top-10 most frequent topics will be return as the result.

7.2 Similar publication venues

For retrieving similar publication venues and years, firstly we need to measure the similarity between two publication venues and years (e.g How similar are "TKDE'14" and "CIKM'15") using the paper titles. One simple way is to firstly tokenize every title of the venues and years, and use bag-of-words model, computing the cosine similarity of the word frequency vectors between them. However, based on our statistical analysis, the dimensionality is too high. It is very inefficient to compute similarity. Besides, the data may also be very noisy, containing many words that only appears once in the whole dataset.

To address the aforementioned problems, we adopted Latent Dirichlet allocation (LDA)[1] to learn the latent topics from the data. Instead of using word frequency vectors, the similarity is defined based on the cosine similarity between

⁷<http://nlp.stanford.edu/software/tagger.shtml>

Table 3: Top-10 most popular research topics of the papers published in 2013

Rank	Topic	Frequency
1	wireless sensor networks	1412
2	case study	1182
3	performance analysis	568
4	cognitive radio networks	468
5	special issue	452
6	performance evaluation	404
7	cloud computing	373
8	genetic algorithm	341
9	empirical study	340
10	neural network	340

the topic distribution between two venues and years. Thus, for a query venue and year, we compute the topic similarities between it and every other venues and years. Then, top-10 similar venues and years are returned as the result.




Figure 2: Graph representation of Latent Dirichlet allocation (LDA)

7.3 Application Experiments

For retrieving popular research topics, we have the results as follows:

For finding similar publication venues and years, we firstly learn the topics as follows:

The results of the queries are given as:

8. CONCLUSIONS

9. REFERENCES

- [1] D. M. Blei, A. Y. Ng, and M. I. Jordan. Latent dirichlet allocation. *the Journal of machine Learning research*, 3:993–1022, 2003.
- [2] M. Ley et al. dblp computer science bibliography dataset. <http://dblp.uni-trier.de/xml>, 2016. last accessed: 18-March-2016.

Table 4: Representative words for the topics leant by LDA

Topic	Words
Geographic Information System	data using based sar spatial sensing analysis remote images radar model land image satellite surface detection gis classification resolution hyperspectral
Robotics	robot control based using robots mobile motion robotic planning human autonomous design navigation vision dynamic sensor multi force tracking visual
Wireless Sensor Networks	networks wireless based network mobile routing ad hoc sensor protocol performance control using traffic efficient scheme qos algorithm 802 multi
Software Technology	software testing analysis based using test engineering development study code case model systems java quality source empirical tool approach program
Social Media	social online media networks community network behavior use internet effects study analysis communities understanding games self role game impact influence

Table 5: Top-10 most similar venues & years to “IEEE Trans. Knowl. Data Eng. 2014”

Rank	Venue	Year	Similarity
1	IEEE Trans. Knowl. Data Eng.	2015	0.9898
2	IEEE Trans. Knowl. Data Eng.	2012	0.9888
3	DaWaK	2015	0.9887
4	IEEE Trans. Knowl. Data Eng.	2013	0.9856
5	LD4IE@ISWC	2014	0.9850
6	CIKM	2006	0.9845
7	WAIM	2013	0.9830
8	CIKM	2007	0.9823
9	FQAS	2009	0.9815
10	CIKM	2004	0.9803